

# Fine Tuning in Practice

## Activate Python Environment and Set API Key

```
dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10$ python -m venv venv

dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10$ 
dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10$ 
dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10$ 
dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10$ 
dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10$ ls
venv
dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10$ . venv/bin/activate
(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10$ 
(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10$ ls
homework1 homework2 homework3 homework4 venv
(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10$ cd homework
-bash: cd: homework: No such file or directory
(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10$ cd homework1
(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10/homework1$ cat ../../week6/homework1/.env
API_KEY=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
ORG_ID=org-A00vXLRbnib0TuY2mcn7PB5A
GOOG_API_KEY=AIzaSyDR9QgS1yJUR1asJ7guelp5fx4LRLRN57U
(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10/homework1$ export OPENAI_API_KEY="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"
(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10/homework1$ |
```

## Create JSONL File

- Create a **JSONL** file, **data.json** with **prompt-completion pairs**.
  - **JSONL format** is **newline-delimited JSON** that is suitable for **training data**.
- We are using **3 examples** in our **JSONL** file, so the **resulting model** will not **perform well**, but we are testing things out and we want to see how it works in practice for now.

```
(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10/homework1$ code .
(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10/homework1$ cat data.json
{"prompt": "When do I have to start the heater?", "completion": "Every day in the morning at 7AM. You should stop it at 2PM."}
{"prompt": "Where is the garage remote control?", "completion": "Next to the yellow door, on the key ring"}
{"prompt": "Is it necessary to program the scent diffuser every day?", "completion": "The scent diffuser is already programmed, you just need to recharge it when its battery is low"}(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10/homework1$ |
```

## Analyze and Prepare Data

```
(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10/homework1$ openai tools fine_tunes.prepare_data -f data.json
Analyzing...
- Your JSON file appears to be in a JSONL format. Your file will be converted to JSONL format
- Your file contains 3 prompt-completion pairs. In general, we recommend having at least a few hundred examples. We've found that performance tends to linearly increase for every doubling of the number of examples
- All prompts end with suffix '?'
- Your data does not contain a common ending at the end of your completions. Having a common ending string appended to the end of the completion makes it clearer to the fine-tuned model where the completion should end. See https://platform.openai.com/docs/guides/fine-tuning/preparing-your-dataset for more detail and examples.
- The completion should start with a whitespace character (' '). This tends to produce better results due to the tokenization we use. See https://platform.openai.com/docs/guides/fine-tuning/preparing-your-dataset for more details

Based on the analysis we will perform the following actions:
- [Necessary] Your format 'JSON' will be converted to 'JSONL'
- [Recommended] Add a suffix ending '\n' to all completions [Y/n]: Y
- [Recommended] Add a whitespace character to the beginning of the completion [Y/n]: Y

Your data will be written to a new JSONL file. Proceed [Y/n]: Y
Wrote modified file to 'data_prepared.jsonl'
Feel free to take a look!

Now use that file when fine-tuning:
> openai api fine_tunes.create -t "data_prepared.jsonl"

After you've fine-tuned a model, remember that your prompt has to end with the indicator string `?` for the model to start generating completions, rather than continuing with the prompt. Make sure to include `stop: ["\n"]` so that the generated texts ends at the expected place.
Once your model starts training, it'll approximately take 2.48 minutes to train a 'curie' model, and less for 'ada' and 'babbage'. Queue will approximately take half an hour per job ahead of you.
(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10/homework1$ cat data_prepared.jsonl
>{"prompt": "When do I have to start the heater?", "completion": " Every day in the morning at 7AM. You should stop it at 2PM\n"}, {"prompt": "Where is the garage remote control?", "completion": " Next to the yellow door, on the key ring\n"}, {"prompt": "Is it necessary to program the scent diffuser every day?", "completion": " The scent diffuser is already programmed, you just need to recharge it when its battery is low\n"}
```

## Fine-Tune the Model

Use the **OpenAI CLI tool** to initiate the **fine-tuning process** using the **prepared JSONL data file**:

```
openai api fine_tunes.create -t "data_prepared.jsonl" -m curie
```

```
(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10/homework1$ openai api fine_tunes.create -t "data_prepared.jsonl" -m curie
Upload progress: 100%
Uploaded file from data_prepared.jsonl: file-y1l1mmjE2X33RVMASmlRz008f
Created fine-tune: ft-UokKswUWeRQ7VLuy0lMWXgRZ
Streaming events until fine-tuning is complete...

([Ctrl-C will interrupt the stream, but not cancel the fine-tune)
[2023-11-19 20:22:34] Created fine-tune: ft-UokKswUWeRQ7VLuy0lMWXgRZ
[2023-11-19 20:23:00] Fine-tune costs $0.00
[2023-11-19 20:23:01] Fine-tune enqueued. Queue number: 0
[2023-11-19 20:23:02] Fine-tune started
```

## Resume Fine-Tuning (if needed)

If the **fine-tuning operation is interrupted**, you can **resume it** using this command:

```
openai api fine_tunes.follow -i <YOUR_FINE_TUNE_JOB_ID>
```

*I skipped this step because my fine tuning was not interrupted.*

## List Fine-Tuned Models

List your **fine-tuned models** using this command:

```
openai api fine_tunes.list
```

```
(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10/homework1$ openai api fine_tunes.list
{
  "object": "list",
  "data": [
    {
      "object": "fine-tune",
      "id": "ft-UokkswUWeRQ7VLuy0lMWXgRZ",
      "hyperparams": {
        "n_epochs": 4,
        "batch_size": 1,
        "prompt_loss_weight": 0.01,
        "learning_rate_multiplier": 0.1
      },
      "organization_id": "org-A00vXLRbnib0TuY2mcN7PB5A",
      "model": "curie",
      "training_files": [
        {
          "object": "file",
          "id": "file-y11mmmjE2X33RVMASmlRz008f",
          "purpose": "fine-tune",
          "filename": "data_prepared.jsonl",
          "bytes": 417,
          "created_at": 1700454153,
          "status": "processed",
          "status_details": null
        }
      ],
      "validation_files": [],
      "result_files": [
        {
          "object": "file",
          "id": "file-BVVRMeQklr0vVWLTz2KR4PPx",
          "purpose": "fine-tune-results",
          "filename": "compiled_results.csv",
          "bytes": 746,
          "created_at": 1700454265,
          "status": "processed",
          "status_details": null
        }
      ],
      "created_at": 1700454154,
      "updated_at": 1700454154
    }
  ]
}
```

## Use the Fine-Tuned Model

Set the **fine-tuned model** and use it for **completions**:

```
export FINE_TUNED_MODEL=<FINE_TUNED_MODEL>
openai api completions.create -m $FINE_TUNED_MODEL -p
<YOUR_PROMPT>
```

```
(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10/homework1$ export FINE_TUNED_MODEL="curie:ft-personal-2023-11-20-04-24-24"
(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10/homework1$ openai api completions.create -m $FINE_TUNED_MODEL -p "When do I need to start my air conditioner?"
When do I need to start my air conditioner?
1. If you just want a cold mist on your face for 1(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10/homework1$ |
```

## Use Fine-Tuned Model in Python

If using **Python**, use the **fine-tuned model** as shown below:

```
import openai
openai.Completion.create(
    model=FINE_TUNED_MODEL,
    prompt=YOUR_PROMPT
    # additional parameters
    # temperature,
    # frequency_penalty,
    # presence_penalty
    # ..etc
)
```

```
(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10/homework1$ cat test.py
import openai
response = openai.Completion.create(
    model="curie:ft-personal-2023-11-20-04-24-24",
    prompt="When do I need to start my air conditioner?"
    # additional parameters
    # temperature,
    # frequency_penalty,
    # presence_penalty
    # ..etc
)
print(response.choices[0].text)(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10/homework1$ python test.py

There are many ways to build this one. Stainless Steel and Bronze
```

## Use Fine-Tuned Model with cURL

Use **cURL** to interact with the **fine-tuned model**:

```
curl https://api.openai.com/v1/completions \
      -H "Authorization: Bearer $OPENAI_API_KEY" \
      -H "Content-Type: application/json" \
      -d '{"prompt": YOUR_PROMPT, "model": FINE_TUNED_MODEL}'
(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10/homework1$ curl https://api.openai.com/v1/completions -H "Authorization: Bearer $OPENAI_API_KEY" -H "Content-Type: application/json" -d '{"prompt": "When do I need to start my air conditioner?", "model": "curie:ft-personal-2023-11-20-04-24-24"}'
{
  "id": "cmpl-8MqZLAIFXIL5Bcfsx5CEiRgjfnpm",
  "object": "text_completion",
  "created": 1700455501,
  "model": "curie:ft-personal-2023-11-20-04-24-24",
  "choices": [
    {
      "text": "\n\nInstall the air conditioner so that the air conditioner outlet is level",
      "index": 0,
      "logprobs": null,
      "finish_reason": "length"
    }
  ],
  "usage": {
    "prompt_tokens": 11,
    "completion_tokens": 16,
    "total_tokens": 27
  }
}
```

## Analyze Fine-Tuned Model

Use this command to analyze the **fine-tuning results**:

```
openai api fine_tunes.results -i <YOUR_FINE_TUNE_JOB_ID>
```

```
(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10/homework1$ openai api fine_tunes.results -i ft-UokKswUWeRQ7VLuy0lMWXgRZ
step,elapsed_tokens,elapsed_examples,training_loss,training_sequence_accuracy,training_token_accuracy
1,33,1,1.7561375378817319,0.0,0.23529411764705882
2,66,2,1.4656687936349773,0.0,0.55
3,91,3,1.682499381750822,0.0,0.36363636363636365
4,124,4,1.3051996972024575,0.0,0.55
5,157,5,1.5102434629015624,0.0,0.29411764705882354
6,182,6,1.5153488444536922,0.0,0.45454545454545453
7,207,7,1.4586372093359627,0.0,0.45454545454545453
8,240,8,1.3332806890644133,0.0,0.35294117647058826
9,273,9,1.0225329246313777,0.0,0.65
10,306,10,1.0040450236957985,0.0,0.65
11,339,11,1.2657160077011214,0.0,0.35294117647058826
12,364,12,1.3393244941780964,0.0,0.45454545454545453
13,389,13,1.3393244941780964,0.0,0.45454545454545453
14,422,14,0.9928822490677702,0.0,0.65
```

## Add Suffix to Fine-Tuned Model Name

You can add a **suffix** to your **fine-tuned model name** using the **suffix parameter**:

```
openai api fine_tunes.create -t data.jsonl -m <engine> --suffix
"my_model_name"
```

```
(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10/homework1$ openai api fine_tunes.create -t data_prepared.jsonl -m davinci --suffix "ruis_model_name"
Found potentially duplicated files with name 'data_prepared.jsonl', purpose 'fine-tune' and size 417 bytes
file-y1lmmjE2X33RVMASmlRz008f
Enter file ID to reuse an already uploaded file, or an empty string to upload this file anyway:
Upload progress: 100%|██████████| 417/417 [00:00<00:00, 639kit/s]
Uploaded file from data_prepared.jsonl: file-GYyyGTЛИWhJFvAltoyMAmBMt
Created fine-tune: ft-6BEJWUE9LD9sroymLt2s4fe5
Streaming events until fine-tuning is complete...

([Ctrl-C will interrupt the stream, but not cancel the fine-tune)
[2023-11-19 20:51:18] Created fine-tune: ft-6BEJWUE9LD9sroymLt2s4fe5
[2023-11-19 20:51:25] Fine-tune costs $0.01
[2023-11-19 20:51:25] Fine-tune enqueued. Queue number: 0
[2023-11-19 20:51:28] Fine-tune started
```

## Delete Fine-Tuned Model

Delete your **fine-tuned model** using these methods:

```
# CLI
openai api models.delete -i <FINE_TUNED_MODEL>

# Python
openai.Model.delete(FINE_TUNED_MODEL)

# cURL
curl -X "DELETE" https://api.openai.com/v1/models/<FINE_TUNED_MODEL>
-H "Authorization: Bearer $OPENAI_API_KEY"
```

```
(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10/homework1$ openai api models.delete -i davinci:ft-personal:ruis-model-name-2023-11-20-04-53-41
{
  "id": "davinci:ft-personal:ruis-model-name-2023-11-20-04-53-41",
  "object": "model",
  "deleted": true
}
(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10/homework1$ openai api models.delete -i "curie:ft-personal-2023-11-20-04-24-24"
{
  "id": "curie:ft-personal-2023-11-20-04-24-24",
  "object": "model",
  "deleted": true
}
(venv) dingrui@LAPTOP-MUAISB9U:/mnt/c/sfbu/cs589/week10/homework1$ |
```